

```
2019-09-21 17:02:10,735 INFO org.apache.hadoop.yarn.client.RMProxy (main): Connecting to
ResourceManager at ip-172-31-42-27.ec2.internal/172.31.42.27:8032
2019-09-21 17:02:13,481 INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat (main):
Total input files to process : 1
2019-09-21 17:02:13,508 INFO com.hadoop.compression.lzo.GPLNativeCodeLoader (main): Loaded
native gpl library
2019-09-21 17:02:13,512 INFO com.hadoop.compression.lzo.LzoCodec (main): Successfully
loaded & initialized native-lzo library [hadoop-lzo rev
7e6c862e89bc8db32c064454a55af74ddff73bae]
2019-09-21 17:02:13,655 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): number of
splits:20
2019-09-21 17:02:13,905 INFO org.apache.hadoop.mapreduce.JobSubmitter (main): Submitting
tokens for job: job_1569085160789_0001
2019-09-21 17:02:14,814 INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl (main):
Submitted application application_1569085160789_0001
2019-09-21 17:02:14,998 INFO org.apache.hadoop.mapreduce.Job (main): The url to track the
job: http://ip-172-31-42-27.ec2.internal:20888/proxy/application_1569085160789_0001/
2019-09-21 17:02:15,005 INFO org.apache.hadoop.mapreduce.Job (main): Running job:
job_1569085160789_0001
2019-09-21 17:02:27,209 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1569085160789_0001 running in uber mode : false
2019-09-21 17:02:27,210 INFO org.apache.hadoop.mapreduce.Job (main): map 0% reduce 0%
2019-09-21 17:02:54,417 INFO org.apache.hadoop.mapreduce.Job (main): map 3% reduce 0%
2019-09-21 17:02:56,429 INFO org.apache.hadoop.mapreduce.Job (main): map 7% reduce 0%
2019-09-21 17:03:06,486 INFO org.apache.hadoop.mapreduce.Job (main): map 12% reduce 0%
2019-09-21 17:03:07,502 INFO org.apache.hadoop.mapreduce.Job (main): map 19% reduce 0%
2019-09-21 17:03:08,507 INFO org.apache.hadoop.mapreduce.Job (main): map 23% reduce 0%
2019-09-21 17:03:09,514 INFO org.apache.hadoop.mapreduce.Job (main): map 28% reduce 0%
2019-09-21 17:03:10,518 INFO org.apache.hadoop.mapreduce.Job (main): map 33% reduce 0%
2019-09-21 17:03:12,536 INFO org.apache.hadoop.mapreduce.Job (main): map 35% reduce 0%
2019-09-21 17:03:13,543 INFO org.apache.hadoop.mapreduce.Job (main): map 41% reduce 0%
2019-09-21 17:03:14,549 INFO org.apache.hadoop.mapreduce.Job (main): map 43% reduce 0%
2019-09-21 17:03:15,554 INFO org.apache.hadoop.mapreduce.Job (main): map 47% reduce 0%
2019-09-21 17:03:16,559 INFO org.apache.hadoop.mapreduce.Job (main): map 51% reduce 0%
2019-09-21 17:03:18,573 INFO org.apache.hadoop.mapreduce.Job (main): map 52% reduce 0%
2019-09-21 17:03:19,583 INFO org.apache.hadoop.mapreduce.Job (main): map 56% reduce 0%
2019-09-21 17:03:20,587 INFO org.apache.hadoop.mapreduce.Job (main): map 58% reduce 0%
2019-09-21 17:03:21,594 INFO org.apache.hadoop.mapreduce.Job (main): map 61% reduce 0%
2019-09-21 17:03:22,600 INFO org.apache.hadoop.mapreduce.Job (main): map 63% reduce 0%
2019-09-21 17:03:31,662 INFO org.apache.hadoop.mapreduce.Job (main): map 67% reduce 0%
2019-09-21 17:03:33,670 INFO org.apache.hadoop.mapreduce.Job (main): map 73% reduce 0%
2019-09-21 17:03:34,674 INFO org.apache.hadoop.mapreduce.Job (main): map 75% reduce 0%
2019-09-21 17:03:35,678 INFO org.apache.hadoop.mapreduce.Job (main): map 78% reduce 0%
2019-09-21 17:03:38,691 INFO org.apache.hadoop.mapreduce.Job (main): map 80% reduce 0%
2019-09-21 17:03:39,697 INFO org.apache.hadoop.mapreduce.Job (main): map 83% reduce 0%
2019-09-21 17:03:40,701 INFO org.apache.hadoop.mapreduce.Job (main): map 85% reduce 0%
2019-09-21 17:03:41,706 INFO org.apache.hadoop.mapreduce.Job (main): map 88% reduce 0%
2019-09-21 17:03:43,715 INFO org.apache.hadoop.mapreduce.Job (main): map 93% reduce 0%
2019-09-21 17:03:44,720 INFO org.apache.hadoop.mapreduce.Job (main): map 98% reduce 0%
2019-09-21 17:03:45,724 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 0%
2019-09-21 17:03:50,748 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 11%
2019-09-21 17:03:57,782 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 22%
2019-09-21 17:03:58,787 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 33%
2019-09-21 17:03:59,790 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 44%
```

```
2019-09-21 17:04:01,797 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 56%
2019-09-21 17:04:02,801 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 77%
2019-09-21 17:04:03,805 INFO org.apache.hadoop.mapreduce.Job (main): map 100% reduce 100%
2019-09-21 17:04:05,820 INFO org.apache.hadoop.mapreduce.Job (main): Job
job_1569085160789_0001 completed successfully
2019-09-21 17:04:05,959 INFO org.apache.hadoop.mapreduce.Job (main): Counters: 56
```

File System Counters

```
FILE: Number of bytes read=62489917
FILE: Number of bytes written=160349028
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=2120
HDFS: Number of bytes written=0
HDFS: Number of read operations=20
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
S3: Number of bytes read=1319447352
S3: Number of bytes written=87522407
S3: Number of read operations=0
S3: Number of large read operations=0
S3: Number of write operations=0
```

Job Counters

```
Killed map tasks=1
Killed reduce tasks=1
Launched map tasks=20
Launched reduce tasks=10
Data-local map tasks=20
Total time spent by all maps in occupied slots (ms)=59154048
Total time spent by all reduces in occupied slots (ms)=15429312
Total time spent by all map tasks (ms)=1232376
Total time spent by all reduce tasks (ms)=160722
Total vcore-milliseconds taken by all map tasks=1232376
Total vcore-milliseconds taken by all reduce tasks=160722
Total megabyte-milliseconds taken by all map tasks=1892929536
Total megabyte-milliseconds taken by all reduce tasks=493737984
```

Map-Reduce Framework

```
Map input records=85331845
Map output records=85331845
Map output bytes=961483442
Map output materialized bytes=92935036
Input split bytes=2120
Combine input records=85331845
Combine output records=15362582
Reduce input groups=6626985
Reduce shuffle bytes=92935036
Reduce input records=15362582
Reduce output records=6626985
Spilled Records=30725164
Shuffled Maps =180
Failed Shuffles=0
Merged Map outputs=180
GC time elapsed (ms)=20694
CPU time spent (ms)=543520
Physical memory (bytes) snapshot=19965620224
```

```
Virtual memory (bytes) snapshot=108274487296
Total committed heap usage (bytes)=17933795328
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=1319447352
File Output Format Counters
  Bytes Written=87522407
```